

The text layer is only trusted where it matches the page

Each of the 27 judgments read twice — once from the PDF's text layer, once by OCR — and the two compared. The four named files decode every letter correctly and every digit wrong.

